

## DW schema and the problem of views

- ▷ DW is multidimensional
- ▷ Schemas: Stars and constellations
- ▷ Typical DW queries, TPC-H and TPC-R benchmarks
- ▷ Views and their materialization

Main references [PJ01, C<sup>+</sup>01, Joh02, Kot02]

# Multidimensional Data Model

Warehouse data is often *thought of* as multidimensional, yet frequently *stored* as relations.

# Multidimensional Data Model

- ▷ Multidimensional Data Model(s)
- ▷ Registrar's Example

# A Conceptual Multidimensional Data Model

Data are

▷ facts, eg “Jill took CS5678”

facts have numeric *measure values* (eg 87%)

▷ dimensions, eg “Jill is a *grad* student”

“CS4002 is an *undergrad* course.”

A *cube* represents facts, by mapping dimension combos to measures.

measure value, for fact "Jill got 87 in CS5678"

|        | Jill | Jack | Al |
|--------|------|------|----|
| CS5678 | 87   |      | 73 |
| EE2222 |      | 25   | 95 |
| CS4002 |      | 89   | 62 |

"students" dimension

"courses" dimension

empty cell, no fact about Jill and CS4002


## More on cubes

Cube generalizes idea of spreadsheet. Can have any number of dimensions.

Another measure for fact “Jill took CS5678” is the 503 times she yawned.

## Registrar's Example

3-d cube, Session  $\times$  Student  $\times$  Course  $\rightarrow$  Grade  
mapping session, student, course to grades.





# Registrar's High-Dimensional Cube

Realistic cubes can have dozens of dimensions.

Registrar:

Session  $\times$  Student  $\times$  Course  $\times$  Section  $\times$  Campus  $\times$  ForCredit  $\rightarrow$  Grade

## Dimensions can have structure

Dimension values can be organized into **containment hierarchies**.

CS5678 is contained in “Grad Courses”

EE2222 and CS4002 are in “U/g Courses”

CS5678 and CS4002 are in “CS Courses”

## Flat dimensions

$\top$  (“top”) means “ALL”. Simplest kind of dimension is “flat”, in that every value is a child of  $\top$ .

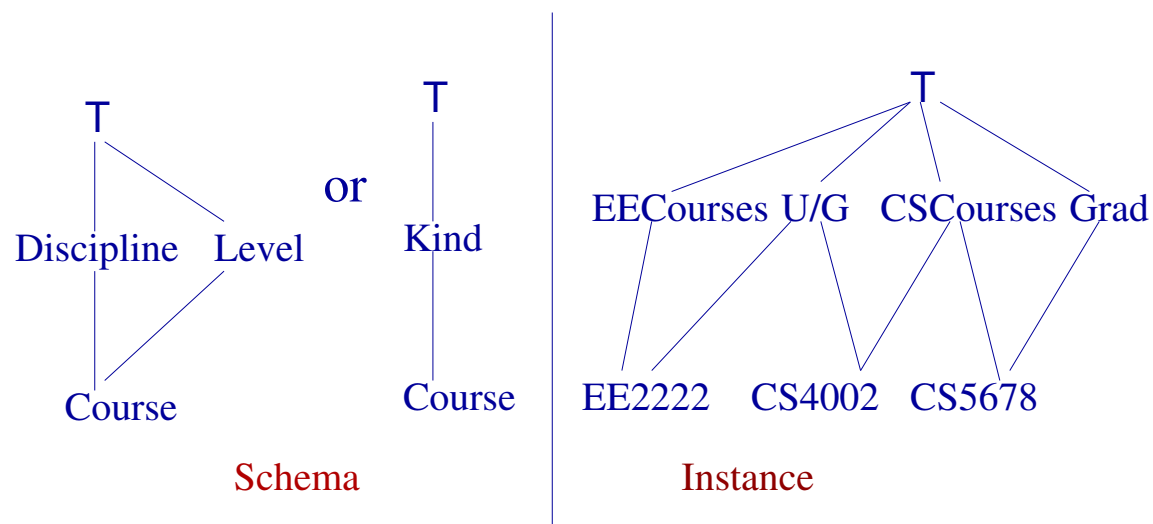
Eg, Campus values are “SJ” and “F”. Or  $\top$ .



Schema (“Type info”) vs Instance (“Values”)

$\top$  overloaded as type and value; we have 3 Campus values!

## Dimensions: can be tricky



How to handle this? Like multiple inheritance.

Some multidimensional models forbid general lattice: demand balanced tree structure in instance.

## Measures are Interesting Too!

In Data Warehousing analyses, we aggregate.

Measure has

- ▷ set of “measure values” eg real numbers  
(report a numeric property of a fact)
- ▷ aggregation operation  
(eg SUM, AVG, MEDIAN)  
→ distributive ones are nicest (SUM).

## Measures are Interesting Too! (part 2)

Best to consider the aggregation operation as “hard-coded” into a cube.

There might be more than one measure to keep track of: SUM and AVG for example.

# Views

Conceptually, views are derived relations computed “on the fly”.

But high-performance database systems often secretly *materialize* (pre-compute, store and keep updated) views for later use. ■

It's not the *analyst's* business whether this is done... but it is a big issue for the DW *designer*.

Crudely, queries should be answered from the smallest summary table possible. ■ Query-optimizer software should take care of this.

## Aggregate Views

What about “CSCourses”, we said it was a dimension value too!

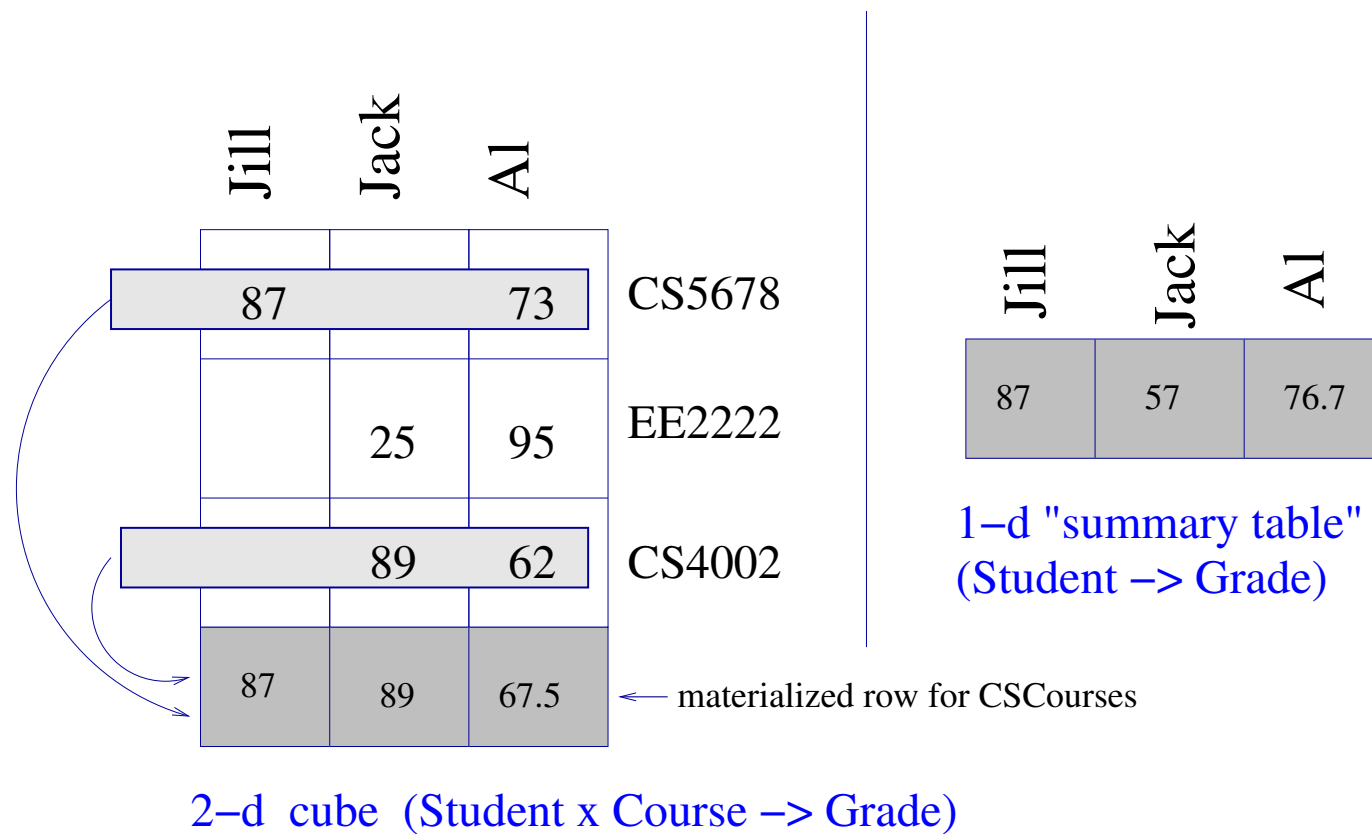
But there are no facts *directly* about anybody’s performance on CSCourses. ■ They are views; compute them by averaging facts involving CS5678 and CS4002.

A materialized view (created by aggregation) is a *summary table*.

Summary tables usually have fewer dimensions than “base cube”.



## Base cube and a summary cube



## Aside: Materialized View Support

Major commercial DBMSs currently support materialized views.

Unfortunately, support has lagged in MySQL and PostgreSQL.■

**MySQL:** no views at all till version 5.■x4! new (&cheap) add-on product to support materialized views and rewrite queries to use them [EBM05].■

**PostgreSQL:** had views a little earlier. “Do It Yourself” materialization is possible via the postgres scripting language and triggers (see Gardner [Gar04]).■Gardner also project lead of *matview Project* for PostgreSQL [mP03].■Project appears inactive.

## Basic Operations

An important part of any data model is “what operations are provided?”. Deferred to a later lecture.

Soon, we describe relational *storage*, and a class of SQL queries.

## Other multidimensional models

Many different multidimensional models.

Some models, eg., Thomsen's "LC" [Tho02] are not as rigid about separating "dimensions" from "measures".

LC distinguishes valid data from "missing" and "meaningless".

## Storing cubes in a relational database

Relational technology is mature, robust, available.

But somehow, we must map cubes into world of tuples and tables.

- ▷ Star Schema (now)■
- ▷ Snowflake Schema (later lecture)■
- ▷ Constellation (now)

# Star Schema

- ▷ One main **fact table**
- ▷ one **dimension table** per dimension.

■  
Fact table stores measure value AND for each dimension, a *foreign key* into the corresponding dimension table.

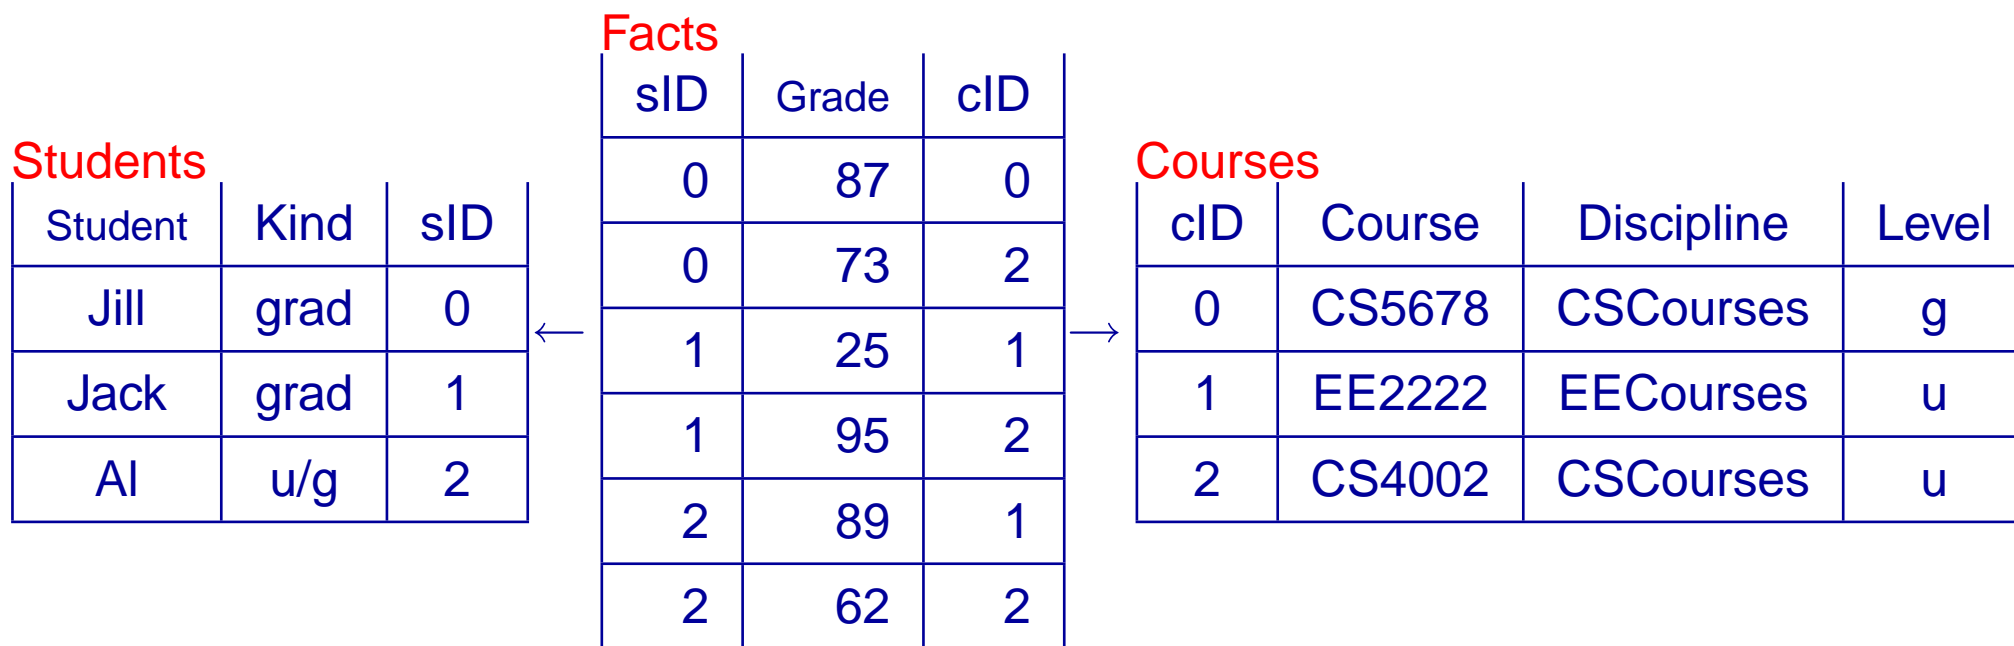
Dimension table stores name of (lowest level) dimensional value *and all* higher-level dimension values that contain it.■

Dimension table's key is a “dummy value” and might be the record-id number.

■

## Star Schema Example

Name comes from the fact table in the centre, with all the dimension tables radiating out from the centre. ■



(no column necessary for  $\top$  in dim tables.)

# Constellation

*A fact constellation:* several fact tables sharing common dimensions.■

Eg, Parking-office fact table

Student  $\times$  Campus  $\rightarrow$  Number of Parking Violations

“Student” dimension, at least, would be same as with registrar’s star



## DW Queries

Aggregation is the tool to “see the forest through the trees”.■

Common form of SQL query is

```
SELECT attrs, aggr-expr...  
FROM fact,dim1,dim2,...WHERE ...  
GROUP BY attrs
```

■

```
SELECT kind, avg(Grade) FROM Facts, Students  
WHERE Facts.sID = Students.sID  
GROUP BY kind
```

■

## DW Query Example, Cont.

```
SELECT kind, avg(Grade) FROM Facts, Students  
WHERE Fact.sID = Students.sID GROUP BY kind
```

is probably implemented by a **star join** between Facts and Students:■

| Student | kind | sID | Grade | cID |
|---------|------|-----|-------|-----|
| Jill    | grad | 0   | 87    | 0   |
| Jill    | grad | 0   | 73    | 2   |
| Jack    | grad | 1   | 25    | 1   |
| Jack    | grad | 1   | 95    | 2   |
| Al      | u/g  | 2   | 89    | 1   |
| Al      | u/g  | 2   | 62    | 2   |

then

| kind | avg(Grade) |
|------|------------|
| grad | 70         |
| u/g  | 75.5       |

## DW Query Example: Using Summary Tables

```
SELECT kind, avg(Grade) FROM Facts,Students  
WHERE Fact.sID = Students.sID GROUP BY kind
```

But suppose that we already had a summary table giving the

average grade for each student:

| Student | kind | sID | Avg(Grade) |
|---------|------|-----|------------|
| Jill    | grad | 0   | 80         |
| Jack    | grad | 1   | 60         |
| Al      | u/g  | 2   | 75.5       |

Then, we might try to average Jack's avg grade with Jill's, to get an average grad grade. *Might* be faster.■

But **avg() is not distributive!** (eg Jack took 1 course, Jill took 10).

## Summary Tables: When faster for queries?

Summary table typically much smaller than the fact table.

When you *can* correctly answer a query from a summary table, speed *might* be helped by smaller data.

... ■ But the main fact tables may have fast index built for it, which the summary table lacks.

Moral: query optimization is tricky.

## Aside: Performance benchmarks for DW systems

The hardware, OS, and database all interact and are crucial for overall DW/OLAP performance. Cannot easily separate issues.

“Transaction Processing Council”: neutral party that evaluates (commercial) DB systems. See <http://www.tpc.org>

■  
Their TPC-**H** benchmark simulates an ad-**h**oc (OLAP-style) environment.

Their TPC-**R** benchmark simulated “canned” queries typical of **r**eport generators. Declared obsolete in 2005.

■

## Aside: Performance for DW systems, 2

TPC-H computes “Composite Queries-per-hour” at various database sizes.

As of 26 Feb 2006, for a 10TiB warehouse, champion performance is Oracle database on a SunFire platform running Solaris 10.

For a 100GiB warehouse, champ is SQL Server 2005 on a HP Proliant running Windows Server 2003.

Source:[http://www.tpc.org/tpch/results/tpch\\_perf\\_results.asp](http://www.tpc.org/tpch/results/tpch_perf_results.asp)

Microsoft/Intel strong on lower end, commercial databases on Linux in the midsize, large multiprocessors running commercial UNIX and Oracle/DB2 at upper end.

## Aside: Factoring price into TPC-H

TPC-H also rates systems by Price/Performance. (System cost/composite queries per hour). Again, affected by database size.

Source: [http://www.tpc.org/tpch/results/tpch\\_price\\_perf\\_results.asp](http://www.tpc.org/tpch/results/tpch_price_perf_results.asp)

10TiB: typically \$6M USD, large multiprocessor. Saw one \$14M.■

100GiB: typically \$100k 4-cpu system■

### Winners:

10TiB, 100GiB champs are the performance champs.■

1TiB, champ is SQL Server 2005 on Bull NovaScale with 16 Itanium-2 CPUs, running Windows 2003. Cost: abt \$400k.■

## Choosing summary tables in DW design

DW designer (or database system) needs to decide which summary tables to use. Issues:

- ▷ identifying candidates ■
- ▷ potential size of each ■
- ▷ potential query benefit of each ■
- ▷ potential cost of keeping each updated ■

Do we know the queries that will be asked?



# Identifying Candidate Summary Tables

Common approach: consider summaries that are

▷ group-bys ■

▷ of some subset ■

▷ of fact-table columns

Candidates belong to the “data cube” [GBLP97].

## Data-Cube Aggregates

Eg, for Facts(sID, cID, Grade) the possible candidates are

1. Select sID, cID, avg(Grade) from Facts group by sID, cID (= Facts!)
2. Select sID, avg(Grade) from Facts group by sID
3. Select cID, avg(Grade) from Facts group by cID
4. Select avg(Grade) from Facts

Note that if we choose to materialize 2 and 4, we can first compute 2.  
Then 4 from 2.

(Likewise, 4 can be computed from 3)

## Sizes of Summary Tables

Various sophisticated techniques can try to predict the size of a summary table. Deferred till next week.

## Benefit of Summary Table

Basically, depends on smallness

... and how often it can be used to answer queries.

Deferred till next week.

## References

- [C<sup>+</sup>01] Surajit Chaudhuri et al. Database technology for decision support systems. *IEEE Computer*, pages 48–55, December 2001.
- [EBM05] EBM Software BV. X4! online at <http://x4.olap4all.com/>, 2005. checked 25 February 2006.
- [Gar04] Jonathan Gardner. Materialized views in PostgreSQL. online,

[http://www.jonathangardner.net/PostgreSQL/materialized\\_views/matviews.html](http://www.jonathangardner.net/PostgreSQL/materialized_views/matviews.html),  
2004. checked 25 February 2006.

- [GBLP97] J. Gray, A. Bosworth, A. Layman, and H. Pirahesh. Data cube: a relational aggregation operator generalizing group-by, cross-tabs and subtotals. Number 1, pages 29–53, 1997.
- [Joh02] Theodore Johnson. Data warehousing. In J. Abello et al., editors, *Handbook of Massive Data Sets*, chapter 19, pages 661–710. Kluwer, 2002.
- [Kot02] Yannis Kotidis. Aggregate view management in data warehouses. In J. Abello et al., editors, *Handbook of Massive Data Sets*, chapter 20, pages 711–741. Kluwer, 2002.
- [mP03] matview Project. Implementation of materialized views for PostgreSQL.

online: <http://gborg.postgresql.org/project/matview/projdisplay.php>, 2003.  
checked 25 February 2006.

- [PJ01] Torben Bach Pedersen and Christian S. Jensen. Multidimensional database technology. *IEEE Computer*, pages 40–46, December 2001.
- [Tho02] Erik Thomsen. *OLAP Solutions*. Wiley, 2nd edition, 2002.